CultureCLIP: Empowering CLIP with Cultural Awareness

Through Synthetic Images and Contextualized Captions

Shayan Salehi

Department of Computer Engineering
Sharif University of Technology

May 2026

Outline

The Challenge of Multimodal Cultural Understanding

Vision-Language Models (VLMs) like CLIP:

- Excel at general cross-modal retrieval and zero-shot classification.
- Capture coarse-grained semantics well (e.g., identifying a "dog" or "deity").
- **Fail at fine-grained cultural nuances.**

The "Cultural Gap"

- Culturally distinct concepts often share overwhelming visual similarities.
- VLMs miss subtle, context-dependent visual cues (e.g., accessories, symbolic colors).
- Consequence: Mismatches and misclassifications.

The Core Question

How can we teach VLMs to capture subtle visual details so they can differentiate between cultural concepts that look similar but mean entirely different things?

Visualizing the CLIP Failure Mode

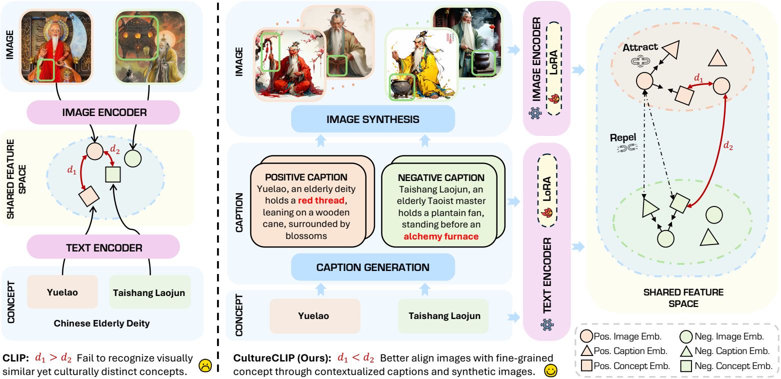

Figure: CLIP's embedding space fails to separate visually similar but culturally distinct concepts (e.g., "Sushi" vs. "Gimbap").

Three Roadblocks to Cultural Alignment

1. **Scarcity of High-Quality Cultural Data**
   Manual curation is expensive and small-scale. Web-scraped data is noisy, culturally misleading, and lacks strict alignment.

2. **CLIP's Preference for Concise Captions**
   CLIP enforces a 77-token limit. Dense, information-rich cultural context disrupts the alignment learned by the model if naively injected.

3. **Absence of Fine-Grained Hard Negatives**
   Standard contrastive learning uses random in-batch negatives. It rarely pairs conceptually similar but visually distinct examples, which are critical for learning cultural boundaries.

Proposed Solution: CulTwin & CultureCLIP

Contribution 1: CulTwin

- A highly curated, synthetic cultural dataset.
- Features **Twin Cards**: pairs of concept-caption-image triplets.
- Triplets are visually similar but culturally distinct, acting as perfect hard negatives.

Contribution 2: CultureCLIP

- A refined contrastive learning framework.
- Jointly aligns concepts, context-rich captions, and synthetic images.
- Parameter-efficient fine-tuning (LoRA) to preserve broad VLM generalizations.

Goal: Achieve state-of-the-art fine-grained cultural recognition without losing general domain capabilities.

CulTwin: A Three-Stage Data Curation Pipeline

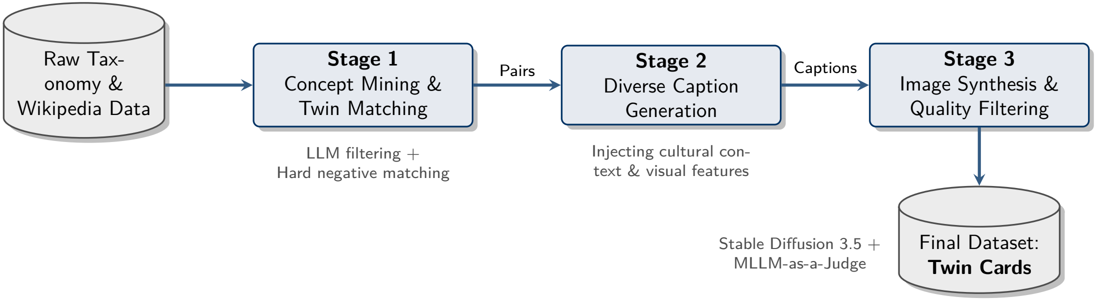

Stage 1 & 2: Concept Mining and Caption Generation

Stage 1: Concept Mining & Twin Matching

- **Taxonomy:** 229 countries across 8 categories (Cuisine, Clothing, Architecture, etc.).
- **LLM Refinement:** We leverage cutting-edge Large Language Models (Qwen2.5-VL) to strictly evaluate cultural relevance.
- **Twin Matching:** The LLM identifies a ”cultural twin” for each concept—a hard negative that is visually similar but culturally distinct.

Stage 2: Diverse Caption Generation

- Generating captions that weave together *both* key visual elements and nuanced cultural background.
- **Diversity Control:** LLM prompts are designed to vary artistic style, lighting, and composition to prevent VLM overfitting on synthetic textures.

Stage 3: Image Synthesis & Quality Filtering

Images generated via Stable Diffusion 3.5 are rigorously evaluated by an **MLLM-as-a-Judge** across three dimensions. Scores below a threshold are discarded.

Table: Image Quality Scores (MLLM-as-a-Judge) & Pass Rates

Taxonomy Category	**Authenticity ↑**	**Consistency ↑**	**Fidelity ↑**	**Pass %**
Cuisine	4.40	3.55	3.75	77.46%
Clothing	4.47	3.86	4.00	75.52%
Animals & Plants	4.35	3.98	4.28	75.63%
Art	4.39	4.03	4.09	72.67%
Architecture	4.22	3.90	4.00	61.32%
Symbol	3.80	3.83	3.98	**85.12%**

Note: Tangible categories (Cuisine/Clothing) yield higher authenticity, while abstract categories (Architecture/Art) face stricter rejection rates.

The Result: "Twin Cards"

Each curated instance in CulTwin is a **Twin Card**: Two opposing triplets of (C, T, I).

Triplet A (Positive)

Concept (C^+): Sushi (Japan)

Caption (T^+): Vinegared rice paired with raw seafood, emphasizing precise cuts and minimalism.

Image (I^+): [Generated Sushi Image]

VS

Triplet B (Hard Negative)

Concept (C^-): Gimbap (Korea)

Caption (T^-): Sesame oil-seasoned rice rolled with cooked meats and pickled vegetables.

Image (I^-): [Generated Gimbap Image]

Impact: The network is forced to learn that the presence of raw fish vs. cooked fillings (visuals) fundamentally alters the cultural concept.

Preliminaries: Contrastive Alignment Evolution

Standard CLIP Objective: Aligns images I and text T by maximizing cosine similarity for matching pairs while pushing apart random in-batch negatives.

$$\mathcal{L}_{\text{CLIP}} = \mathcal{L}_{\text{I2T}} + \mathcal{L}_{\text{T2I}}$$

NegCLIP & TripletCLIP: Introduced explicit hard negative texts T^- (via perturbation) and images I^- to refine decision boundaries.

$$\mathcal{L}_{\text{TripletCLIP}} = \mathcal{L}_{\text{NegCLIP}}(I^+, T^+, T^-) + \mathcal{L}_{\text{NegCLIP}}(I^-, T^-, T^+)$$

The Missing Piece

Existing methods focus purely on image-text alignment. They lack an abstract semantic anchor (the **Concept**) to bind visually nuanced captions and images to their high-level cultural identity.

CultureCLIP: Architecture & Alignment

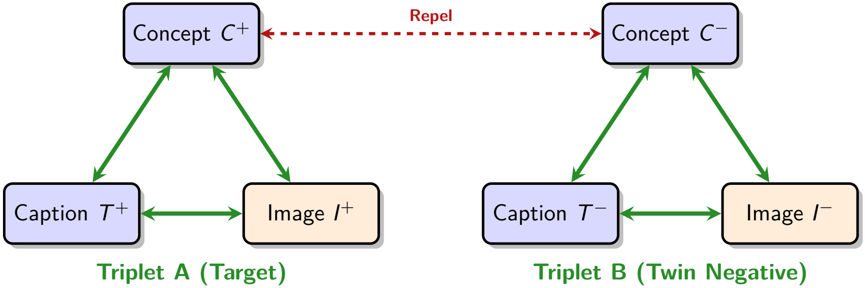

Like a magnetic field: Attract components within a cultural triplet; repel elements from the culturally opposing twin.

CultureCLIP Objective Function

CultureCLIP extends TripletCLIP by treating the abstract **Concept (C)** and detailed **Caption (T)** as dual text anchors sharing the same encoder.

The overall objective balances concept-level and caption-level losses:

$$\mathcal{L}_{\text{CultureCLIP}} = \lambda_c \cdot \mathcal{L}_{\text{concept}} + \lambda_t \cdot \mathcal{L}_{\text{caption}}$$

Where the caption loss enforces text-image alignment against the twin:

$$\mathcal{L}_{\text{caption}} = \mathcal{L}_{\text{NegCLIP}}(I^+, T^+, T^-) + \mathcal{L}_{\text{NegCLIP}}(I^-, T^-, T^+)$$

And the concept loss grounds the image to the abstract cultural identity:

$$\mathcal{L}_{\text{concept}} = \mathcal{L}_{\text{NegCLIP}}(I^+, C^+, C^-) + \mathcal{L}_{\text{NegCLIP}}(I^-, C^-, C^+)$$

Experimental Setup

Training Specifications

- **Base Model:** CLIP (ViT-B/32 image encoder).
- **Method:** Parameter-Efficient Fine-Tuning using **LoRA** (Rank 4).
- **Motivation:** Full fine-tuning disrupts general knowledge representations. LoRA adapts efficiently while preserving pre-trained general space.

Evaluation Benchmarks

- **Culture-Specific:** GlobalRG-G, GlobalRG-R, CROPE.
  (Adapted into statement-ranking tasks evaluating semantic fine-grain accuracy).
- **Culture-Agnostic:** MS COCO, Flickr30k.
  (Testing Recall@5 to ensure standard retrieval abilities aren't forgotten).

Main Results: Overcoming the Cultural Gap

Performance on Culture-Specific and Culture-Agnostic Tasks

Methods	Culture-Specific (Accuracy %)			Culture-Agnostic (R@5)	
	GlobalRG-G	GlobalRG-R	CROPE	MS COCO	Flickr30k
CLIP (Baseline)	63.98	78.22	74.69	65.40	89.00
CLIP++ (Naive FT)	46.05	49.98	73.62	28.80	50.80
NegCLIP++ (Ours)	66.32	78.41	**79.25**	65.50	89.20
CultureCLIP (Ours)	**69.47**	**78.60**	78.84	**66.30**	**89.30**

Key Takeaways:

- **+5.49%** improvement over base CLIP on GlobalRG-G.
- Naive fine-tuning (CLIP++) destroys performance (*Catastrophic Forgetting*).
- CultureCLIP improves cultural awareness *and* general retrieval simultaneously.

Ablation: The Role of Concept vs. Caption

Which contributes more to fine-grained alignment?

Configuration	λ (Cap / Con)	GlobalRG-G	GlobalRG-R	MS COCO
Caption-only (w/ neg)	1.0 / –	66.27	77.53	65.50
Concept-only (w/ neg)	– / 1.0	65.83	77.29	65.60
CultureCLIP (Optimal)	**0.3 / 0.7**	**69.47**	**78.60**	**66.10**

Insight: Concepts act as powerful, abstract semantic anchors. Weighting the concept objective higher ($\lambda_c = 0.7$) allows the model to distinguish subtle cultural boundaries better than relying on dense text alone.

Ablation: LoRA and Data Quality

Configuration	Quality Filtered	LoRA Rank	GlobalRG-G	CROPE
Full FT Baseline	$\times$	–	46.95	73.62
Full FT + QF	✓	–	43.64	74.33
LoRA Baseline	$\times$	4	<u>69.47</u>	78.84
CultureCLIP (+QF + LoRA)	✓	**4**	**69.67**	**79.21**

Insights:

- **LoRA is strictly necessary** to prevent catastrophic forgetting. Full parameter updates on domain-specific data destroy pre-trained manifolds.
- **Quality filtering works:** Training on 73.8k high-quality filtered samples outperforms the full 100k unfiltered set.

Current Limitations

- **Highly Abstract Reasoning:** CultureCLIP still struggles when visual distinctions are purely stylistic or symbolic rather than physically concrete.
- **Evaluation Bottleneck:** Relying on MLLMs (like Qwen2.5-VL) as judges introduces potential model biases and restricts true interpretability.
- **The Synthetic Domain Gap:** While highly diverse, CulTwin is fully synthetic. A distributional shift exists between generated imagery and real-world, in-the-wild cultural photographs.

Conclusion & Takeaways

1. **CulTwin Dataset:** A novel, scalable pipeline for generating synthetic cultural data, leveraging cutting-edge LLMs and Diffusion models to create hard "Twin Card" negatives.
2. **CultureCLIP Architecture:** A refined contrastive learning objective that explicitly anchors detailed visual/textual information to high-level cultural concepts.
3. **State-of-the-Art Results:** Demonstrated significant improvements on culture-specific benchmarks (+5.49%) while simultaneously enhancing general zero-shot retrieval capabilities.

Thank You!

References I

- Radford, A., et al. (2021). Learning Transferable Visual Models From Natural Language Supervision. *ICML*.
- Yuksekgonul, M., et al. (2022). When and Why Vision-Language Models Behave like Bags-of-Words, and What to Do About It?. *ICLR*.
- Patel, et al. (2024). TripletCLIP: Improving Fine-grained Concept Understanding.
- Bai, et al. (2025). Qwen2.5-VL Technical Report.
- Rombach, R., et al. (2022). High-Resolution Image Synthesis with Latent Diffusion Models. *CVPR*.
- Hu, E. J., et al. (2021). LoRA: Low-Rank Adaptation of Large Language Models. *ICLR*.